

Week 5.2 Backpropagation in Practice

Unrolling Parameters

- unroll a matrix into a big long vector

Example

$$s_1 = 10, s_2 = 10, s_3 = 1$$

$$\Theta^{(1)} \in \mathbb{R}^{10 \times 11}, \Theta^{(2)} \in \mathbb{R}^{10 \times 11}, \Theta^{(3)} \in \mathbb{R}^{1 \times 11}$$

$$D^{(1)} \in \mathbb{R}^{10 \times 11}, D^{(2)} \in \mathbb{R}^{10 \times 11}, D^{(3)} \in \mathbb{R}^{1 \times 11}$$


$$\text{thetaVec} = [\text{Theta1}(:); \text{Theta2}(:); \text{Theta3}(:)];$$

$$\text{DVec} = [\text{D1}(:); \text{D2}(:); \text{D3}(:)];$$

$$\text{Theta1} = \text{reshape}(\text{thetaVec}(1:110), 10, 11);$$

$$\text{Theta2} = \text{reshape}(\text{thetaVec}(111:220), 10, 11);$$

$$\text{Theta3} = \text{reshape}(\text{thetaVec}(221:231), 1, 11);$$



Gradient Checking

- two-side difference:

$$\frac{d}{d\theta} J(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon}$$

- one-side difference:

$$\frac{d}{d\theta} J(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta)}{\epsilon}$$

- check $\text{gradApprox} \approx \text{Dvec}$

Random Initialization (Symmetry breaking)

- Initialize $\Theta_{ij}^{(l)}$ to a random value in $[-\epsilon, \epsilon]$

Putting it Together

Training a Neural Network

1. Randomly initialize the weights
2. Implement forward propagation to get $h_{\Theta}(x(i))$ for any $x(i)$

3. Implement the cost function
4. Implement backpropagation to compute partial derivatives
5. Use gradient checking to confirm that your backpropagation works. Then disable gradient checking.
6. Use gradient descent or a built-in optimization function to minimize the cost function with the weights in θ